

# Orchestration Proposal Presentation

Robert Hensing   Robin Kuipers   Jelle Postma

January 5, 2016

# Orchestration?

# Orchestration?

Wikipedia: Orchestration describes the automated arrangement, coordination, and management of complex computer systems, middleware and services.

# Taxonomy

- ▶ Arrangement
- ▶ Coordination
- ▶ Management

# Taxonomy

- ▶ Arrangement
  - ▶ Allocating machines
  - ▶ Deploying applications
- ▶ Coordination
- ▶ Management

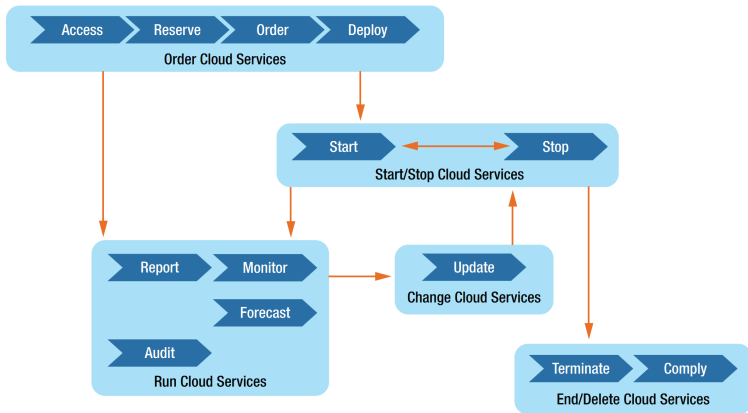
# Taxonomy

- ▶ Arrangement
- ▶ Coordination
  - ▶ Peer discovery
  - ▶ Coordination
- ▶ Management

# Taxonomy

- ▶ Arrangement
- ▶ Coordination
- ▶ Management
  - ▶ Start/stop
  - ▶ Monitoring
  - ▶ Configuration

# Taxonomy



– Open Data Center Alliance



# Existing products ...

... from the hadoop community

**ZooKeeper** A coordination service for distributed applications

**Ambari** Web UI for provisioning, managing and monitoring other hadoop products

**Chukwa** A monitoring / data collection system piggybacking on the storage and compute systems in hadoop

# ZooKeeper

- ▶ Very simple interface
- ▶ Maintenance, coordination, management, consensus. . .
- ▶ Synchronisation for distributed systems
- ▶ Aims to implement functionality that is hard to implement but easy to use

# Ambari

- ▶ Web-based interface
- ▶ Installation, management and monitoring of Hadoop clusters
- ▶ Easily integratable with its RESTful APIs
- ▶ Very wide range of features

# Chukwa

- ▶ Data collection
- ▶ Mostly used for monitoring and analyzing systems
- ▶ Builds on Storage and Computation systems
- ▶ Very scalable and robust

# Choosing between the three

- ▶ Chukwa
- ▶ Ambari
- ▶ Zookeeper

# Choosing between the three

- ▶ Chukwa
  - ▶ Too heavily based on other systems
  - ▶ Therefore an interesting project, but not the scope that we are looking for
- ▶ Ambari
- ▶ Zookeeper

# Choosing between the three

- ▶ Chukwa
- ▶ Ambari
  - ▶ Too large; would have to pick subset of functionality
  - ▶ Choice would again be very dependant on work of other groups
- ▶ Zookeeper

# Choosing between the three

- ▶ Chukwa
- ▶ Ambari
- ▶ Zookeeper
  - ▶ Focusses on concurrent programming
  - ▶ Has a core which we can start implementing, and continue from there
  - ▶ Great for a small project like this



# Choice of project

## TimeKeeper

- ▶ A distributed data store for service orchestration
- ▶ Focus on consistency, availability, scalability
- ▶ Not unlike ZooKeeper
- ▶ Central component of orchestration
- ▶ Greater concurrency
- ▶ Good scope

# Project

- ▶ Deviation from sequential use: for sequential use, we have roughly `Data.Map` with change notifications. Requirement: availability, scalability.
- ▶ We will implement the core logic of an in memory consistent/available data store. Stages, at least three, for example:
  - ▶ Serve from one process
  - ▶ Serve through multiple processes: master/worker
  - ▶ Serve more efficiently with multiple processes: implement sharding (how to deal with consistency?)

# Experimentation

- ▶ Measure req/s for write-intensive vs. read-intensive workloads
- ▶ Measure req/s for single point of synchronisation vs. sharded
- ▶ Measure effect of simulated intercontinental latency on throughput



Questions?