

Orchestration Proposal Presentation

Robert Hensing Robin Kuipers Jelle Postma

January 5, 2016

Orchestration?

Orchestration?

Wikipedia: Orchestration describes the automated arrangement, coordination, and management of complex computer systems, middleware and services.

Taxonomy

Wikipedia level research :)

Taxonomy

- ▶ Arrangement
- ▶ Coordination
- ▶ Management

Taxonomy

- ▶ Arrangement
 - ▶ Allocating machines
 - ▶ Deploying applications
- ▶ Coordination
- ▶ Management

Taxonomy

- ▶ Arrangement
- ▶ Coordination
 - ▶ Peer discovery in distributed app?
- ▶ Management

Taxonomy

- ▶ Arrangement
- ▶ Coordination
- ▶ Management
 - ▶ Monitoring
 - ▶ Configuration

Existing products ...

... from the hadoop community

ZooKeeper A coordination service for distributed applications

Ambari Web UI for provisioning, managing and monitoring other hadoop products

Chukwa A monitoring / data collection system piggybacking on the storage and compute systems in hadoop

ZooKeeper

- ▶ Very simple interface
- ▶ Maintenance, coordination, management, consensus. . .
- ▶ Synchronisation for distributed systems
- ▶ Aims to implement functionality that is hard to implement but easy to use

Ambari

- ▶ Web-based interface
- ▶ Installation, management and monitoring of Hadoop clusters
- ▶ Easily integratable with its RESTful APIs
- ▶ Very wide range of features

Chukwa

- ▶ Data collection
- ▶ Mostly used for monitoring and analyzing systems
- ▶ Builds on Storage and Computation systems
- ▶ Very scalable and robust

Choosing between the three

- ▶ Chukwa
- ▶ Ambari
- ▶ Zookeeper

Choosing between the three

- ▶ Chukwa
 - ▶ Too heavily based on other systems
 - ▶ Therefore an interesting project, but not the scope that we are looking for
- ▶ Ambari
- ▶ Zookeeper

Choosing between the three

- ▶ Chukwa
- ▶ Ambari
 - ▶ Too large; would have to pick subset of functionality
 - ▶ Choice would again be very dependant on work of other groups
- ▶ Zookeeper

Choosing between the three

- ▶ Chukwa
- ▶ Ambari
- ▶ Zookeeper
 - ▶ Focusses on concurrent programming
 - ▶ Has a core which we can start implementing, and continue from there
 - ▶ Great for a small project like this

Choice of project

TimeKeeper

- ▶ A distributed data store for cloud management
- ▶ Focus on consistency, availability, scalability
- ▶ Not unlike ZooKeeper

Project

- ▶ Deviation from sequential use: for sequential use, we have roughly @Data.Map@ with change notifications. Requirement: availability, scalability.
- ▶ We will implement the core logic of an in memory consistent/available data store. Stages, at least three, for example:
 - ▶ Serve from one process
 - ▶ Serve through multiple processes: master/worker
 - ▶ Serve more efficiently with multiple processes: implement sharding (how to deal with consistency?)

Experimentation

- ▶ Measure req/s for write-intensive vs. read-intensive workloads
- ▶ Measure req/s for single point of synchronisation vs. sharded
- ▶ Measure effect of simulated intercontinental latency on throughput



Questions?